


Performance, optimization and why it matters

Marco Scheurer

Lausanne, Switzerland / Shanghai, China



`marco@sente.ch`

WeChat: `phink0`

Performance

- Why?
- What?
- How?
- When?


Why optimizing performance?

- Better user experience

Optimizing performance is...

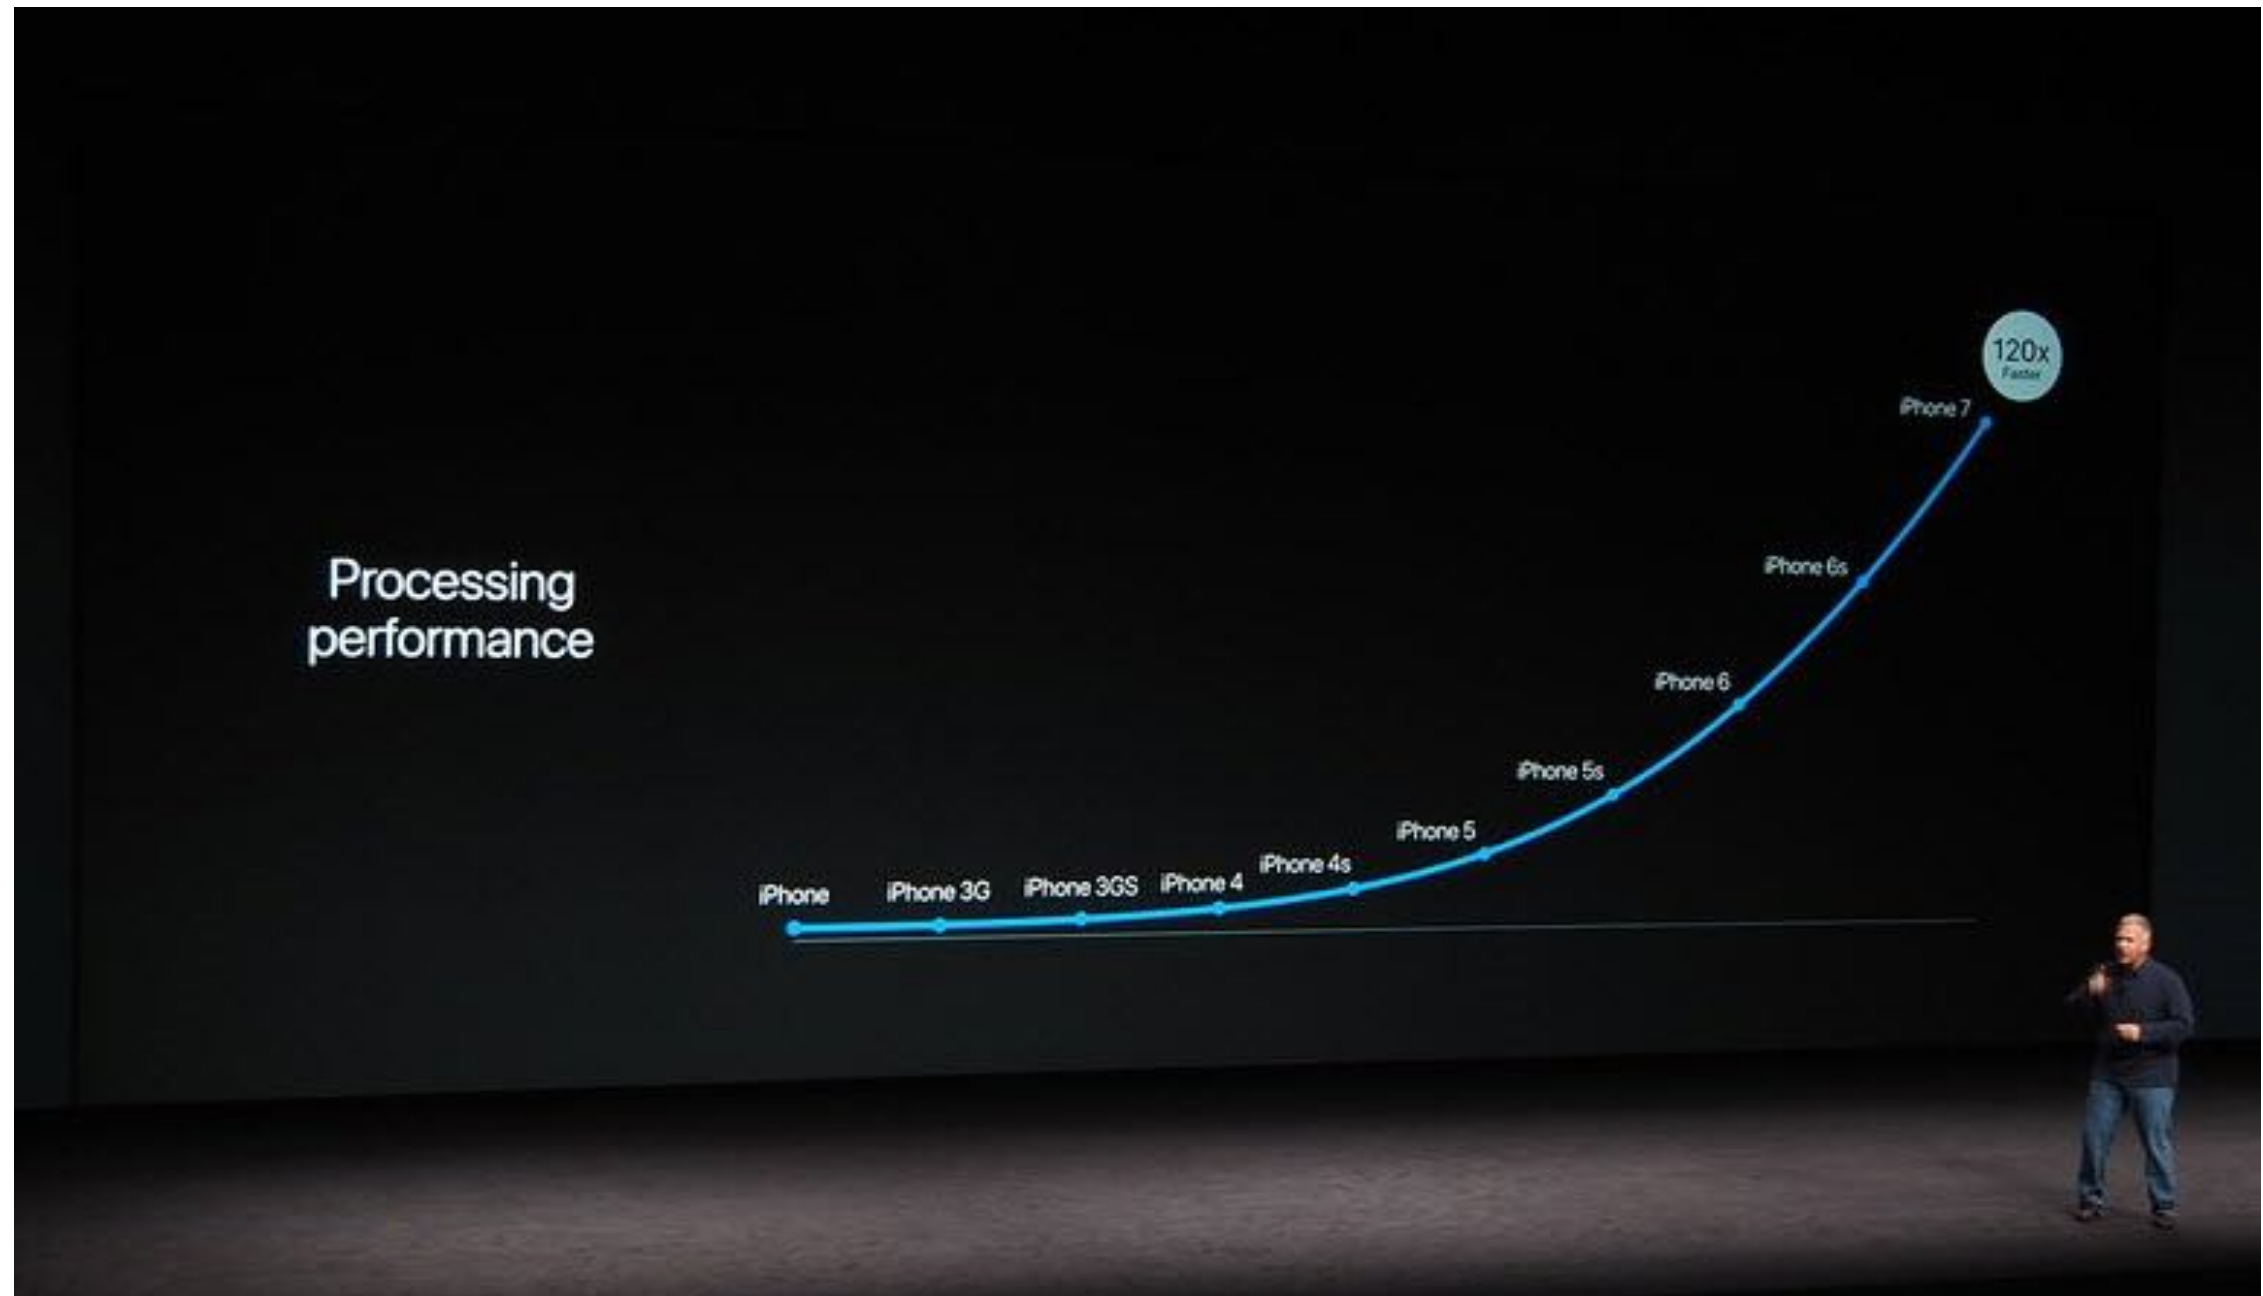
- Risky
“If it ain’t broke, don’t fix it”
- Time consuming
“I don’t have time for performance work”
- Complicated
“I don’t know where to start or what to do”
- Better solved with faster hardware
“Hardware is cheap, programmers are expensive”

Hardware



iPhone 5S, 2013

Hardware



iPhone 7, 2016

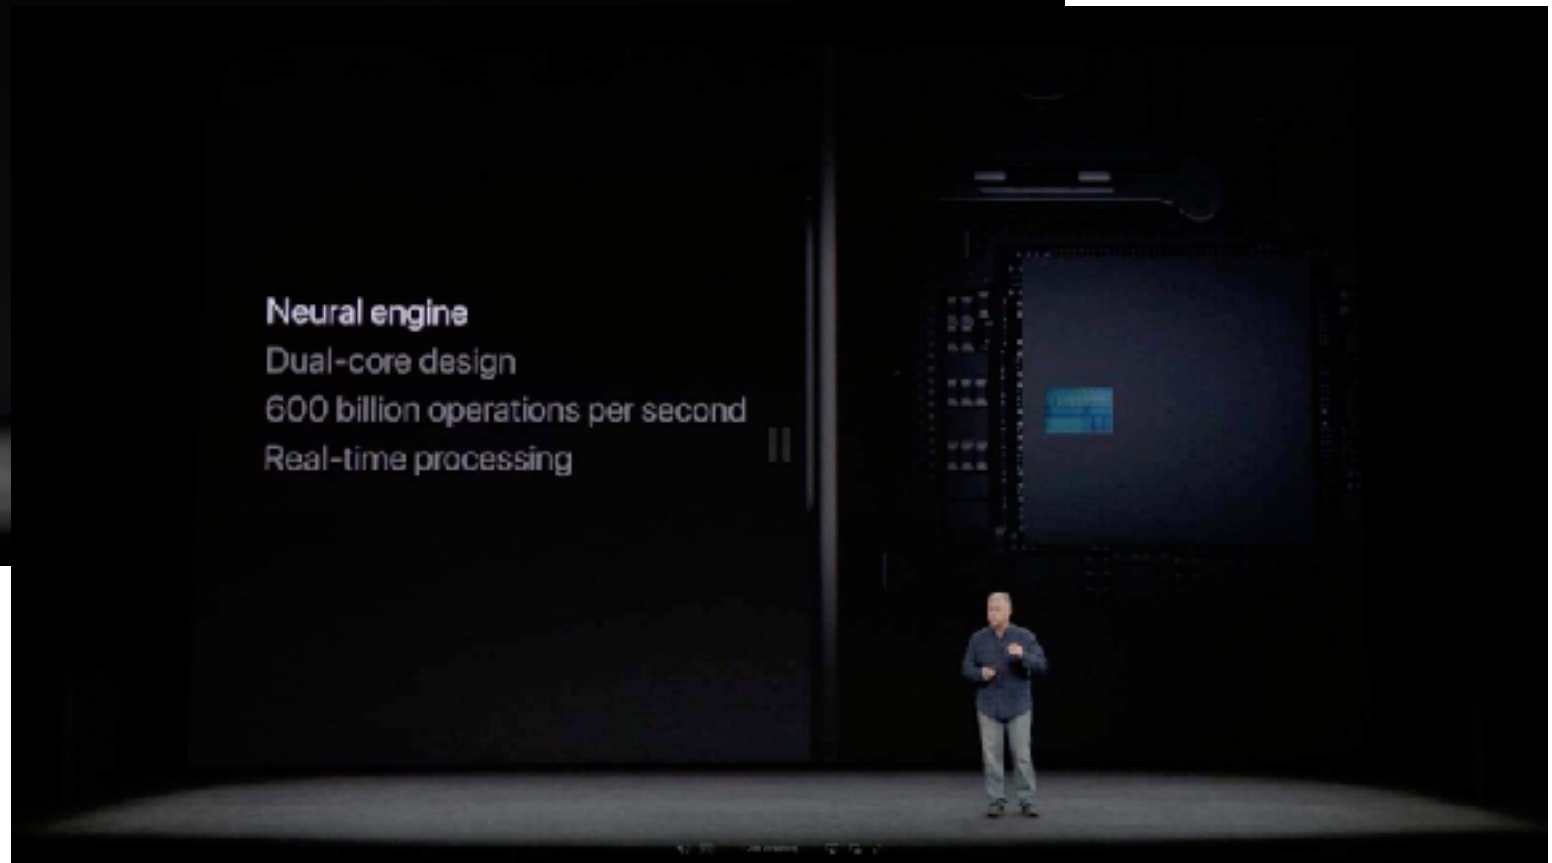
Hardware



Two performance cores
25% faster than A10

Four high-efficiency cores
70% faster than A10

Neural engine
Dual-core design
600 billion operations per second
Real-time processing



iPhone 8, X, 2017

Hardware

- iPhone in 2017 ~200x faster than in 2007 but:
 - Screen with 25x more pixels
 - Augmented reality
 - 240 fps 1080p video capture
 - Animated emojis
 - etc.

Not everybody has an A11


Rank	Device	Usage Share [1]
1	iPhone 7	14.57%
2	iPhone 6S	11.34%
3	iPhone 7 Plus	9.14%
4	iPhone 6	8.26%
5	iPhone X	6.09%
6	iPhone 8	5.02%
7	iPhone 8 Plus	4.88%
8	iPhone SE	4.50%
9	iPhone 6S Plus	4.18%
10	iPad Air 2	4.06%
11	iPhone 5S	3.65%
12	iPad 5	3.54%
13	iPad Mini 2	2.98%
14	iPad Air	2.77%

<https://data.apptelligent.com/ios/devices/>

Device	Device Usage
iPhone 6	13.6%
iPhone 6s	11.0%
iPhone 7	8.9%
iPhone 5s	6.1%
iPhone 8	5.0%
iPhone 7+	4.6%
iPhone se	4.2%
iPhone X	4.2%
iPad Air 2	4.1%
iPhone 8+	4.0%
iPhone 6+	3.6%
iPad Air	3.4%
iPhone 6s+	3.2%
iPad 4G	2.9%
iPod touch 4G	2.7%
iPad 2G	1.9%
iPad Mini 2	1.8%
iPhone 5	1.8%

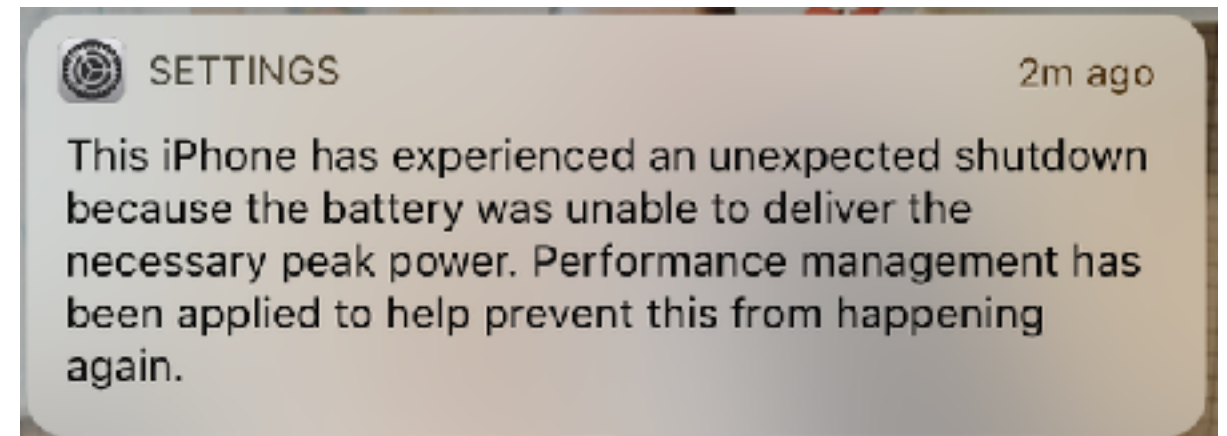
<https://david-smith.org/iosversionstats/>

Battery



- 1400 mAh to 2700 mAh
- Almost no progress, related to device size
- (some Android phones have 5000 mAh)
- iPhone X: 2716 mAh = 10.35 Wh (3.8 V)

Battery



iPhone throttling: New class action takes total to 60 but how many will Apple face?

~ Power consumption (mW)

CPU	500 - 2000
Display	400 - 1000
Cell radio	800
WiFi	200
Bluetooth	40
GPS	200
Microphone	100
Gyroscope	100
Accelerometer	20

<https://qnovos.com/understanding-power-usage-in-a-smartphone/>
https://www.usenix.org/legacy/event/atc10/tech/full_papers/Carroll.pdf

1000 mW on iPhone X drains the battery in 10 hours

Better user experience

1. No crash (including killed by the watchdog)
2. Responsive application
3. Good battery life

Activity Watchdog

- The amount of time the watchdog gives you is not formally documented
- Was more an issue in 2007 than today
- Still happens...

Launch	20 s
Resume	10 s
Suspend	10 s
Quit	5 s
Operation	10 min

Incident Identifier: AD5E7187-C852-4873-8717-89150CCCCF1E
CrashReporter Key: aad64d5368f33ab8bd00f10c7e7414b7067d35ec
Hardware Model: iPhone6,2
Process: **WeChat** [4240]
Path: /private/var/containers/Bundle/Application/
B2CB8F37-61B9-43D3-8F29-C0FBD485E2D3/WeChat.app/WeChat
Identifier: com.tencent.xin
Version: 6.6.2.34 (6.6.2)
Code Type: ARM-64 (Native)
Role: Foreground
Parent Process: launchd [1]
Coalition: com.tencent.xin [1154]
Date/Time: 2018-04-02 12:23:36.3672 +0800
Launch Time: 2018-04-02 12:20:10.5635 +0800
OS Version: iPhone OS 11.2.6 (15D100)
Baseband Version: 8.30.01
Report Version: 104
Exception Type: EXC_CRASH (SIGKILL)
Exception Codes: 0x0000000000000000, 0x0000000000000000
Exception Note: EXC_CORPSE_NOTIFY

Termination Reason: Namespace SPRINGBOARD, Code 0x8badf00d

Termination Description: SPRINGBOARD, process-exit watchdog transgression:

com.tencent.xin exhausted real (wall clock) **time allowance of 5.00**

seconds | | ProcessVisibility: Background | ProcessState: Running |

WatchdogEvent: process-exit | WatchdogVisibility: Background |

WatchdogCPUStatistics: (| "Elapsed total CPU time (seconds): 2.770 (user 2.770,
system 0.000), 28% CPU", | "Elapsed application CPU time (seconds): 0.176, 2%
CPU" |)

Code 0xc00010ff

- “Killed by the operating system in response to a thermal event”
- Another reason to make apps run efficiently

Where to start? Measuring Performance

- Guessing often does not work
- Measure - Change - Measure
- Use measuring tools


Measuring Tools

- NSLog / print
- Instruments
- Simulator
- measure()

NSLog/print

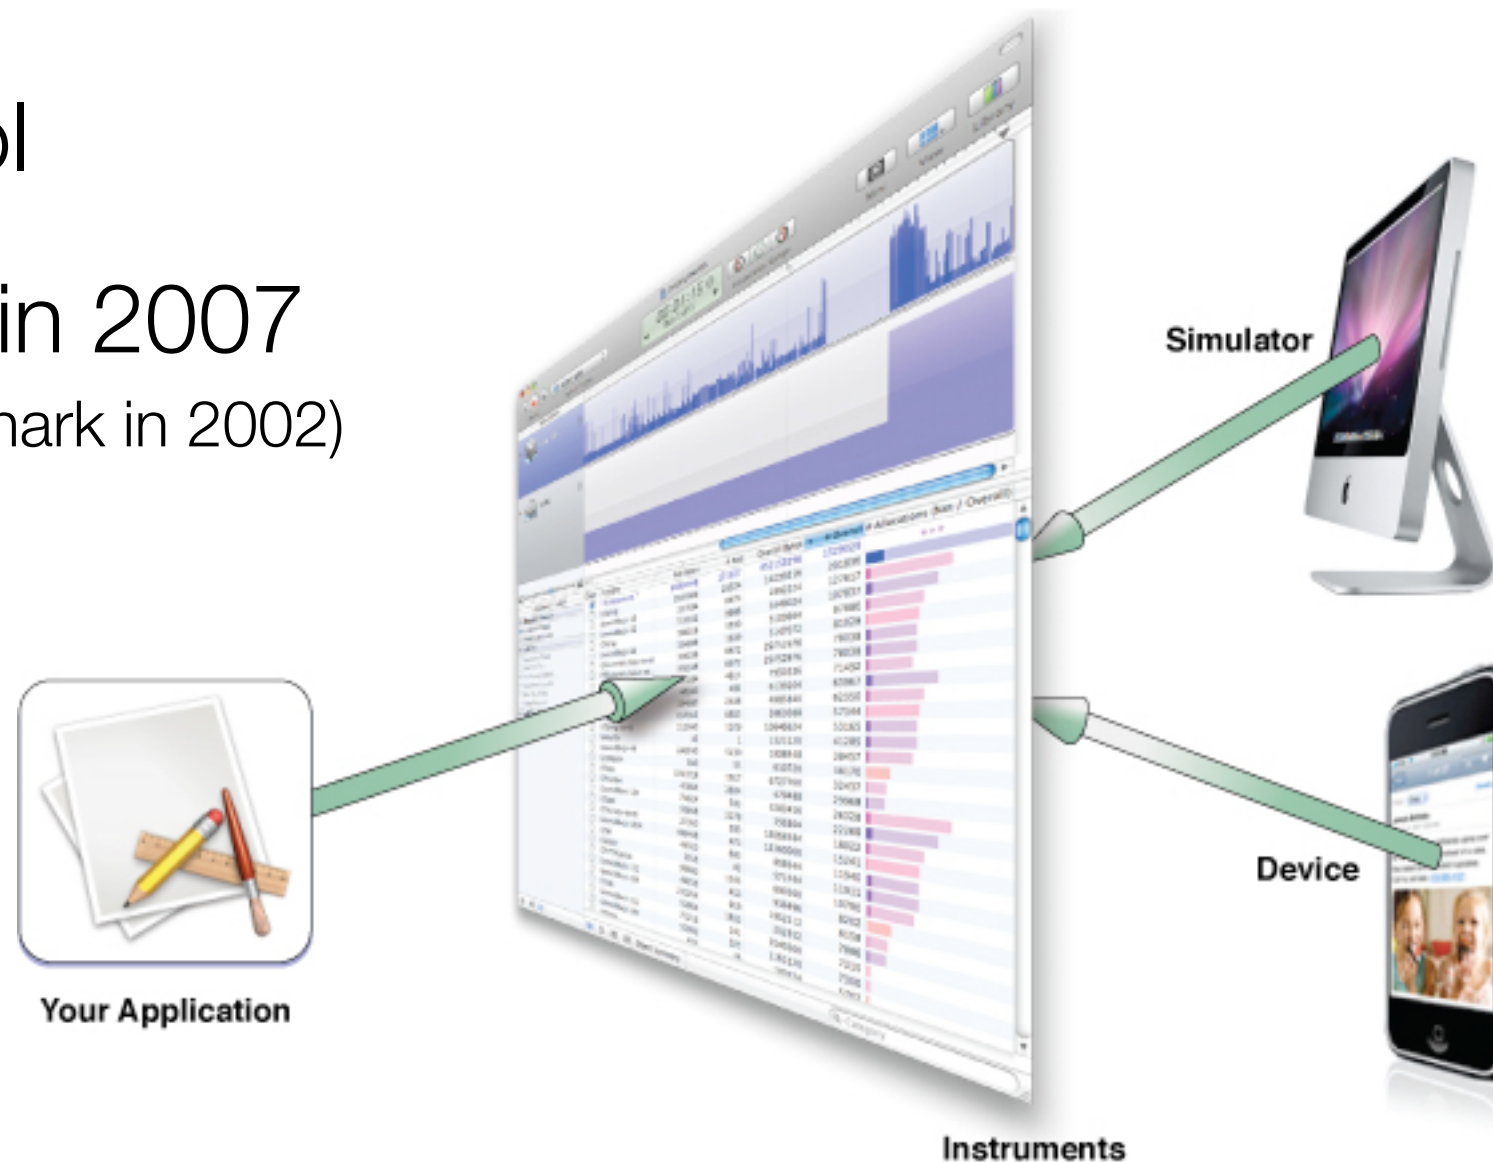
```
NSDate *startTime = [NSDate date];  
...  
NSLog(@"Elapsed time: %.3f", -[startTime timeIntervalSinceNow]);
```


```
double start = CFAbsoluteTimeGetCurrent();  
...  
double elapsedTime = CFAbsoluteTimeGetCurrent() - start;  
NSLog(@"Elapsed time: %.3f", elapsedTime);
```



Instruments

- Profiling tool
- Introduced in 2007
(as Shikari and Shark in 2002)





Instruments



Blank



Activity Monitor



Allocations



Cocoa Layout



Core Animation



Core Data



Counters



Energy Log



File Activity



Leaks



Metal System
Trace



Network



SceneKit



System Trace



System Usage



Time Profiler



Zombies

Demo



Simulator

- Software simulator \neq hardware emulator

Simulator	iPhone	5S	X
i386	ARM	ARM	ARM
2.5 GHz	400 MHz	1.3 GHz	2.4 GHz
Unlimited	128 MB	1GB	3GB

- Measure and test on oldest supported device (5S)
- Simulator OK to check memory usage

XCTestCase: measure()

70 ● Objective C

71 ✓ – (void)testPerformanceExample {

72 ✓ [self measureBlock:^(

73 [self loopContiguous];

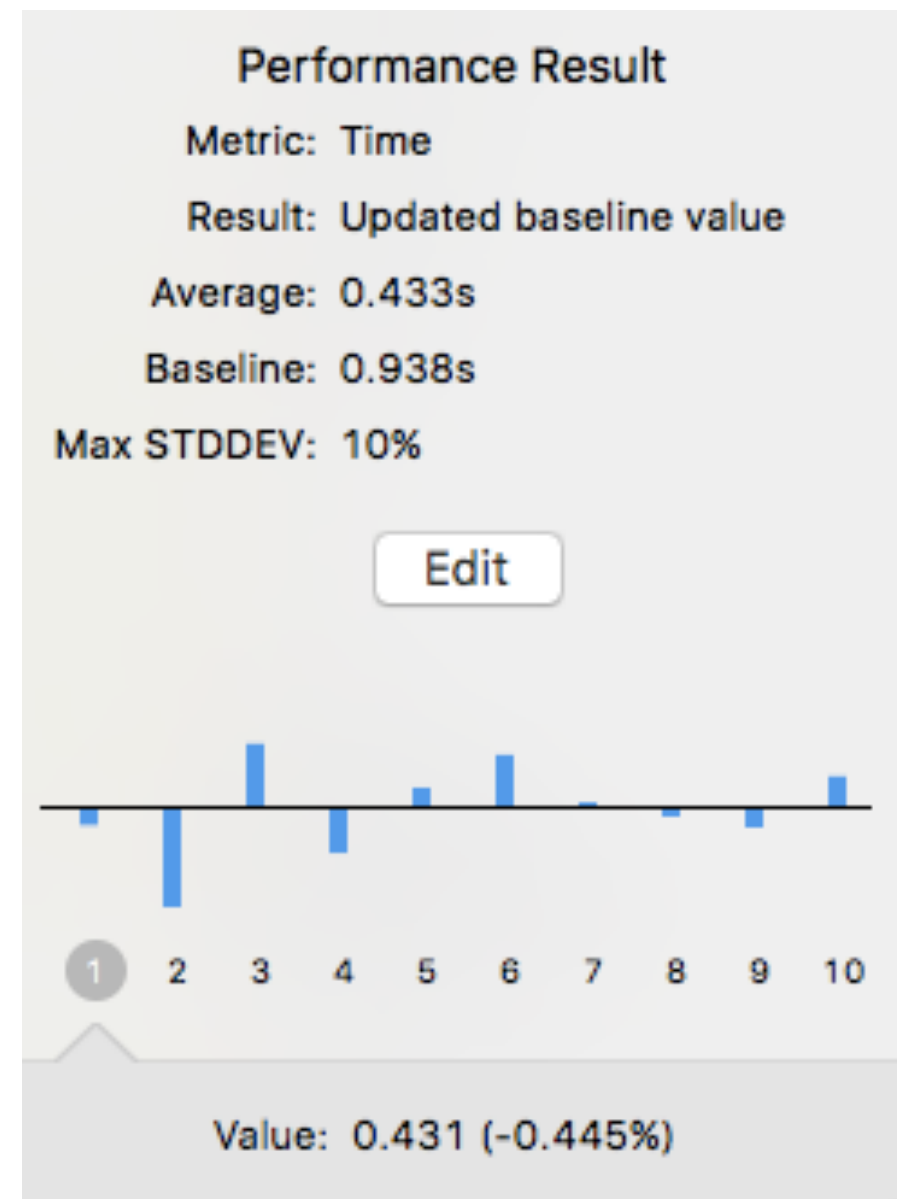
74 }];

75 }

✓ Time: 0.433 sec (54% better)

● Swift

```
func testPerformanceExample() {  
    self.measure {  
    }  
}
```



measure ()

- Set or change the baseline after each run
- Also startMeasuring() / stopMeasuring()
- Match the baseline using device types
- Many tests will slow the suite: run separately
- Demo

1. No crash (including kill by the watchdog)
2. Responsive application
3. Good battery life

1. No Crash





Memory Watchdog

- Watches memory pressure
- Issues low memory warning
- Instant termination of application



Memory

- Swap: no
- Virtual memory: yes
- Memory usage also impacts time



Memory Usage

- Avoid
 - spikes
 - leaks
 - abandoned memory
 - zombies



Memory Spikes

- Be smart with autorelease
- Process large quantities of data in batches
 - Nested autorelease pools
- Demo



Memory Leaks

- Allocated memory that is inaccessible
- Unbalanced retain/release
- Forget to release property's original value
- Leaks Instrument examines heap for leaked memory, identify allocation
- Xcode: Build and Analyze



Abandoned Memory

- Wasted memory
- Accessible but never used
- Memory should not grow without bounds
- Detect with Allocation Instrument



Zombies

- Messages send to deallocated object
- Detect with Zombies instrument
- Demo

1. No crash (including kill by the watchdog)
2. Responsive application
3. Good battery life

2. Responsive applications





Time

- Faster code
 - makes applications more responsive
 - uses less power



Time

- C, ARM code generation
- Objective C
- Frameworks
- File system
- Animation
- Swift



C and ARM Code Generation

- ARM v6 (2G, 3G)
- ARM v7 (3GS, 4)
- ARM 64 (5 ... X)
- -Onone for Debug
- -Os vs -O3 vs -Ofast for Release
(and Profiling and Performance testing)

Code Generation

```
#define MATRIX_SIZE 5000
static double array[MATRIX_SIZE][MATRIX_SIZE];

double sum = 0.0;
for (int j = 0; j < MATRIX_SIZE; j++) {
    for (int i = 0; i < MATRIX_SIZE; i++) {
        array[i][j] = i;
        sum += array[i][j];
    }
}
```

0.55 s

Code Generation

```
#define MATRIX_SIZE 5000  
static double array[MATRIX_SIZE][MATRIX_SIZE];
```

```
double sum = 0.0;  
for (int j = 0; j < MATRIX_SIZE; j++) {  
    for (int i = 0; i < MATRIX_SIZE; i++) {  
        array[i][j] = i;  
        sum += array[i][j];  
    }  
}
```

0.55 s

```
double sum = 0.0;  
for (int i = 0; i < MATRIX_SIZE; i++) {  
    for (int j = 0; j < MATRIX_SIZE; j++) {  
        array[i][j] = i;  
        sum += array[i][j];  
    }  
}
```

0.28 s

● Speedup: ~2x



Objective C

- MYTH: Dynamic dispatch in Objective C is slow
- objc_msgSend() is very, VERY FAST
- ...and very, VERY USEFUL
- IMP caching to remove overhead:
 - Loop with known receiver
 - Rarely useful

objc_msgSend()

```
#define COUNT 100 * 1000 * 1000
```

```
- (uint64_t) withMessageSend
```

```
{  
    NSNumber *n = [NSNumber numberWithInt:42];  
    uint64_t r = 0;  
    for (int i = 0; i < COUNT; i++) {  
        r += [n intValue];  
    }  
    return r;  
}
```

1.37 s

```
- (uint64_t) withIMPCaching
```

```
{  
    NSNumber *n = [NSNumber numberWithInt:42];  
    SEL s = @selector(intValue);  
    int (*f) (id, SEL, ...) = (int (*)(id, SEL, ...))[n methodForSelector:s];  
    uint64_t r = 0;  
    for (int i = 0; i < COUNT; i++) {  
        r += (int)f(n,s);  
    }  
    return r;  
}
```

0.92 s

- Speedup 1.48x but...
- > 200 millions objc_msgSend() / s on iPhone 5S



Frameworks

- Foundation
- Grand Central Dispatch
- Accelerate



Foundation

- NSMutableArray, NSMutableString
 - insert $O(N)$, $O(1)$ at begin and end
- NSMutableDictionary, NSMutableSet
 - lookup $O(1)$ with good hash function
- NSMutableIndexSet
 - Store integers without NSNumber
- NSMutableOrderedSet, NSCountedSet...



Foundation

```
[array indexesOfObjectsWithOptions:NSEnumerationConcurrent  
    passingTest:^(id obj, NSUInteger idx, BOOL *stop) {  
        return [obj test:value];  
    }  
]
```

- Demo



Grand Central Dispatch

- Concurrent programming
- Multicore or single core
- Easier and lighter than threads
- Move long processes off the main thread
- `dispatch_async()`
- `dispatch_apply()` to parallelize loops



Grand Central Dispatch

```
[resultLabel setText:@""];
[progressView startAnimating];

dispatch_queue_t queue = dispatch_get_global_queue(QOS_CLASS_USER_INITIATED, 0);
dispatch_queue_t main_queue = dispatch_get_main_queue();

dispatch_async (queue, ^{
    float result = [self longDotWithAccelerate];
    dispatch_async(main_queue, ^{
        [progressView stopAnimating];
        [resultLabel setText:[NSString stringWithFormat:@"%f", result]];
    });
});
```



Accelerate

- ARM SIMD architecture: NEON
 - 16 8-bit operations at the same time
 - NEON can decode MP3 on 10MHz CPU in real time
- A7 (iPhone 5S)
 - 2 cores
 - 31 general purpose 64-bit registers
 - 32 128-bit vector register
- A11 (iPhone 8, iPhone X)
 - 6 cores
 - Neural Engine, 600 billions ops/s



Accelerate

- 2600 APIs for hardware accelerated math
- vDSP (Digital Signal Processing)
 - Fourier transforms, ...
- BLAS (Basic Linear Algebra Subprograms)
 - Matrix product, ...
- LAPACK (Linear algebra)
 - Solving system of linear equations
- vImage, vForce, BNNS, Compression...

<https://developer.apple.com/wwdc17/711>



Accelerate

```
float dotProduct = 0.0;
for (int i = 0; i < VECTOR_SIZE; i++) {
    dotProduct += a[i] * b[i];
}
```

```
float dotProduct = 0.0;
vDSP_dotpr(a, 1, b, 1, &dotProduct, VECTOR_SIZE);
```

- Makes using NEON and parallel computing easy



Accelerate

```
float dotProduct = 0.0;  
for (int i = 0; i < VECTOR_SIZE; i++) {  
    dotProduct += a[i] * b[i];  
}
```

0.958 s

```
float dotProduct = 0.0;  
vDSP_dotpr(a, 1, b, 1, &dotProduct, VECTOR_SIZE);
```

0.023 s

- Vector size 1024
- 100 000 times on iPhone 5S
- Speedup: 40x



Frameworks

- Demo



File system

- Access to Flash memory is relatively slow
- Speed vary a lot from device to device
- For large files use memory mapped files

`NSDataReadingOptions.DataReadingMappedAlways`

- Long I/O off the main thread



File System

- I/O of small amount of data:
 - UserDefaults
 - Property Lists
 - JSON
 - NSKeyedArchiver
- I/O of large amount of data:
 - Incremental formats
 - Databases (CoreData)



Property Lists

- Binary Format 2-3x faster than XML

- NSPropertyListSerialization

- Slow:

- [NSArray writeToFile:atomically:]
 - [NSDictionary writeToFile:atomically:]
 - [NSString writeToFile:atomically:encoding:error:]

- Fast:

```
NSData *data = [NSPropertyListSerialization dataWithPropertyList:dictionary
                                                    format:NSPropertyListBinaryFormat_v1_0
                                                    options:0
                                                    error:NULL];
[data writeToFile:path atomically:YES];
```

NSJSONSerialization

- Fast
- Portable
- Convenient

```
NSError *error = nil;  
NSArray *json = [NSJSONSerialization JSONObjectWithData:data  
                options: NSJSONReadingMutableContainers  
                error: &error];
```



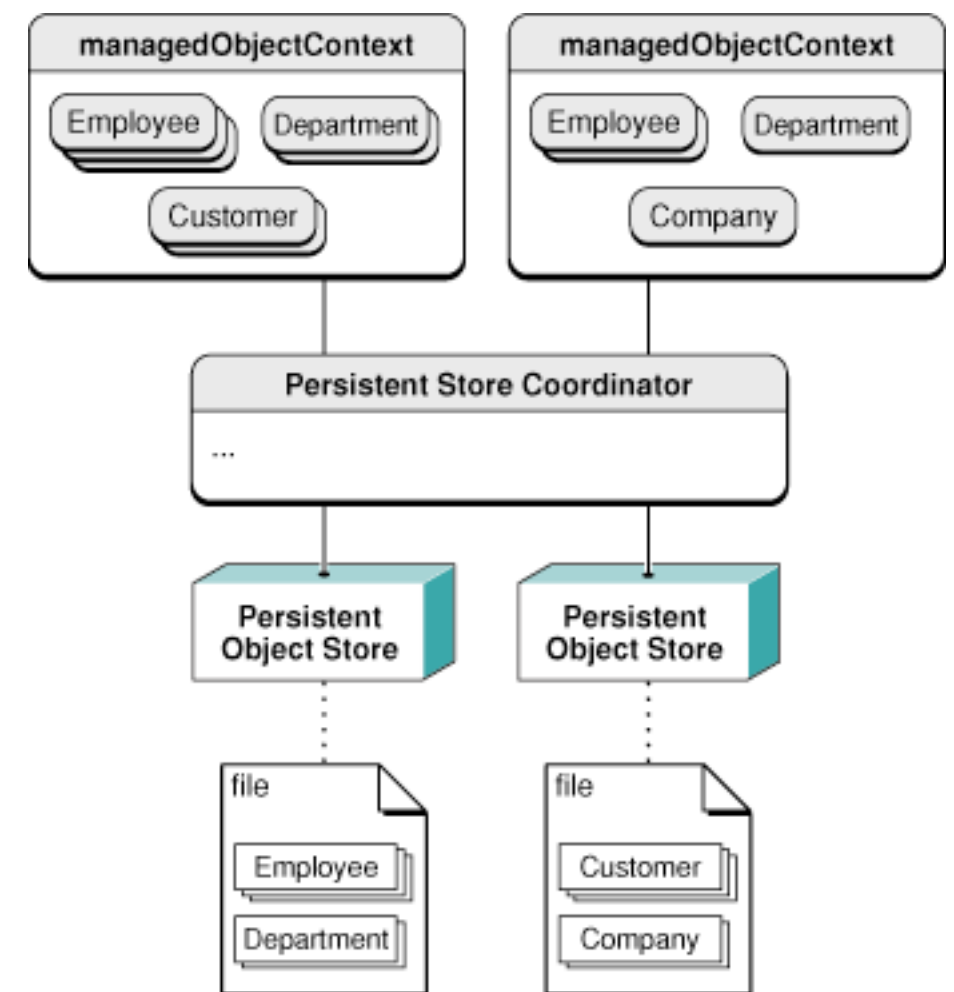
NSKeyedArchiver

- For all objects supporting NSCodering
- Supports versioning
- Not an incremental file format
- Heavy, slow
- Use NSJSONSerialization for “property list”-like values



Database

- Use CoreData... or not
 - fmdb
 - SQLite.swift
 - ... SQL wrappers sometime too much overhead
- Do not import large quantity of data...
- Add object stores instead





Database Search

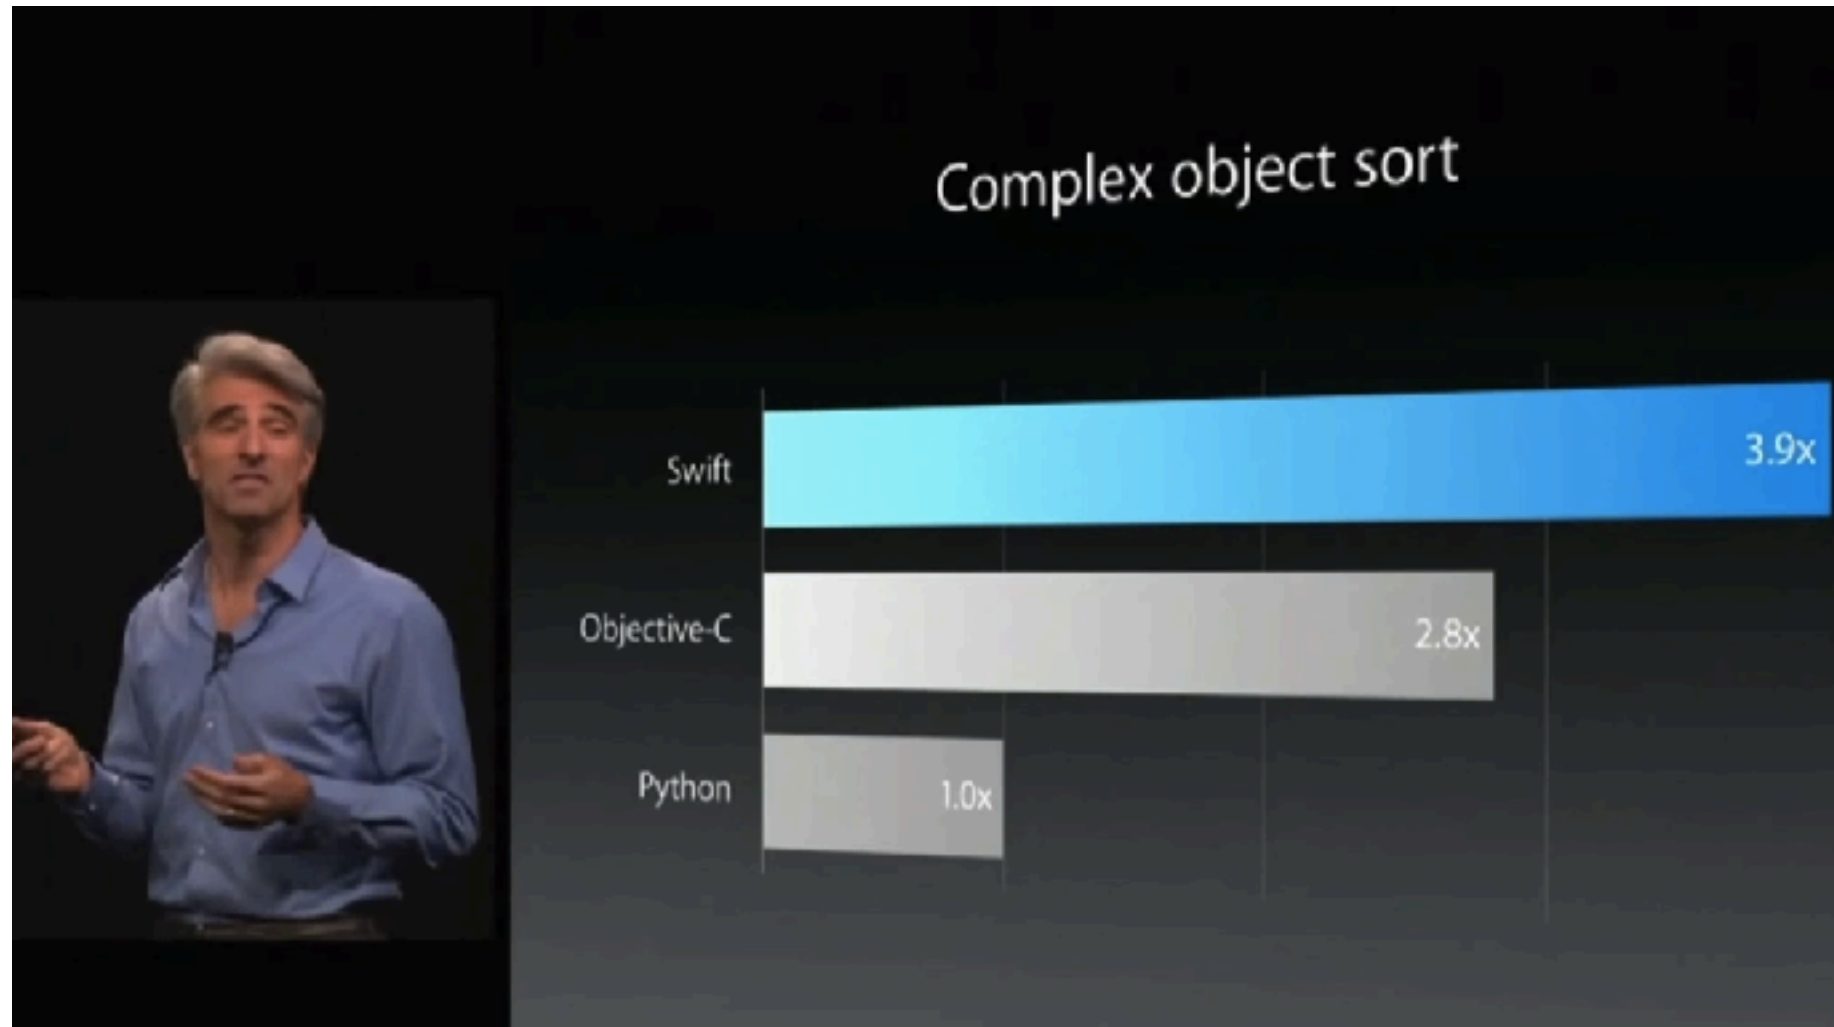
- Unicode string comparisons are expensive
- Use derived attribute (without accents)
- Prefer prefix searching
- Use \leq and $<$ instead of BEGINSWITH



Animation

- Scrolling at 60 fps
 - Reuse identifier
 - Reuse formatters
 - Opaque views
 - UINib

Swift



- Who in the world compares speed with Python?

Swift

Complex object sort



- 1.4 times faster would be nice
- May be true on “complex object sort”
- Not so sure for the real world

Swift Compilation

▼ Swift Compiler - Code Generation

Setting	PerformanceSwift
Disable Safety Checks	No ↕
Exclusive Access to Memory	Full Enforcement (Run-time Checks in Debug Builds Only) ↕
▼ Optimization Level	<Multiple values> ↕
Debug	None [-Onone] ↕
Release	Fast, Whole Module Optimization [-O -whole-module-optimization] ↕
Swift 3 @objc Inference	Default ↕

- When optimizing:
 - Release, not Debug
 - -O -whole-module-optimization
 - Disable runtime safety checks

Inlining

- Automatic inlining by the compiler
- Inline all the things with `@inline(__always)`
- Force inlining with `@_transparent`
- Public interface `@_inlineable`

Source Optimizations

- struct vs class
- final class
- private
- enum vs string

struct vs class

- stack vs heap
- value vs reference
- copy on write
- dynamic dispatch
- reference counting

More Source Optimizations

- map/reduce vs for loop
- Generics
- Protocols
- inout
- Memory allocation
- Unchecked operators (&+)

map vs for loop

```
var array = [Float](repeating: 0, count: 10_000_000)

func useMap () {
    var output = array.map({ (element) -> Float in
        return element * 5
    })
}

func useForLoop () {
    var output = [Float]()
    output.reserveCapacity(array.count)
    for element in array {
        output.append(element * 5)
    }
}
```

map vs for loop

```
var array = [Float](repeating: 0, count: 10_000_000)

func useMap () {
    var output = array.map({ (element) -> Float in
        return element * 5
    })
}

func useForLoop () {
    var output = [Float]()
    output.reserveCapacity(array.count)
    for element in array {
        output.append(element * 5)
    }
}
```

0.10 s

0.08 s

- Speedup: 1.25x
- Not much but still surprising

Generics

- Optimizer will try to specialize the code if the generic definition is visible
- Put generics declaration in module where used
- Or use `-whole-module-optimization`

Protocols

- struct can conform to protocols
- Benefits over class get lost
- references to Protocol Witness Table, Value Witness Table
- 3-words or less inline otherwise on heap

Apple, *Understanding Swift Performance, WWDC 2016*
<https://developer.apple.com/videos/play/wwdc2016/416/>

inout

- Parameters passed to functions:
 1. Function is called: value of argument is copied.
 2. Body of the function: copy is modified.
 3. Function returns: copy's value assigned to original argument.
- *Can* be optimized as a call by reference

Allocation

```
func doSomething () {  
    var s = MyStruct ()  
    DispatchQueue (label:"q").async {  
        s.perform ()  
    }  
}
```

● Heap or Stack ?

Allocation

```
func doSomething () {  
    var s = MyStruct ()  
    DispatchQueue (label:"q").async {  
        s.perform ()  
    }  
}
```

● Heap !

Swift

- Language, compiler and optimizer keep changing
- struct, class, protocol, generics...
 - complicated?
 - low level premature optimization?
- Compiler is slow

1. No crash (including kill by the watchdog)
2. Responsive application
3. Good battery life

3. Battery Life





Energy

- 4G, WiFi, Bluetooth, GPS: 2000 mW
- CPU + GPU: 1000 mW
- Screen: 400 mW



Radios and sensors

- Network
- CoreLocation
- CoreMotion



Radios

- $4G > 3G > 2G > \text{WiFi}$ (LTE = $\sim 5\times$ WiFi)
- 4G / 3G requires radio in high/moderate power state after last packet is sent or received
- LTE: 11 seconds, WiFi: 210 ms
- Low signal uses more power
- Jumping between 4G and 3G/2G drains power
- 2G less power but for a longer time




Network

- Bandwidth usage impacts time and power
- Limit number of connections
- Do not poll
- Use compact data formats
 - (CSV vs JSON vs XML)
 - Compress



CPU

- Do not poll, use events
- Accelerate: more power, less energy





CoreLocation

- Use minimum required accuracy

GPS	kCLLocationAccuracyBest
GPS	kCLLocationAccuracyNearestTenMeters
WiFi	kCLLocationAccuracyHundredMeters
Cell / WiFi	kCLLocationAccuracyKilometer



CoreLocation

- `distanceFilter`
 - how often you receive location changed notifications
 - default: all changes, many events, high CPU usage
- `stopUpdatingLocation`
 - when good enough accuracy, switch GPS off.

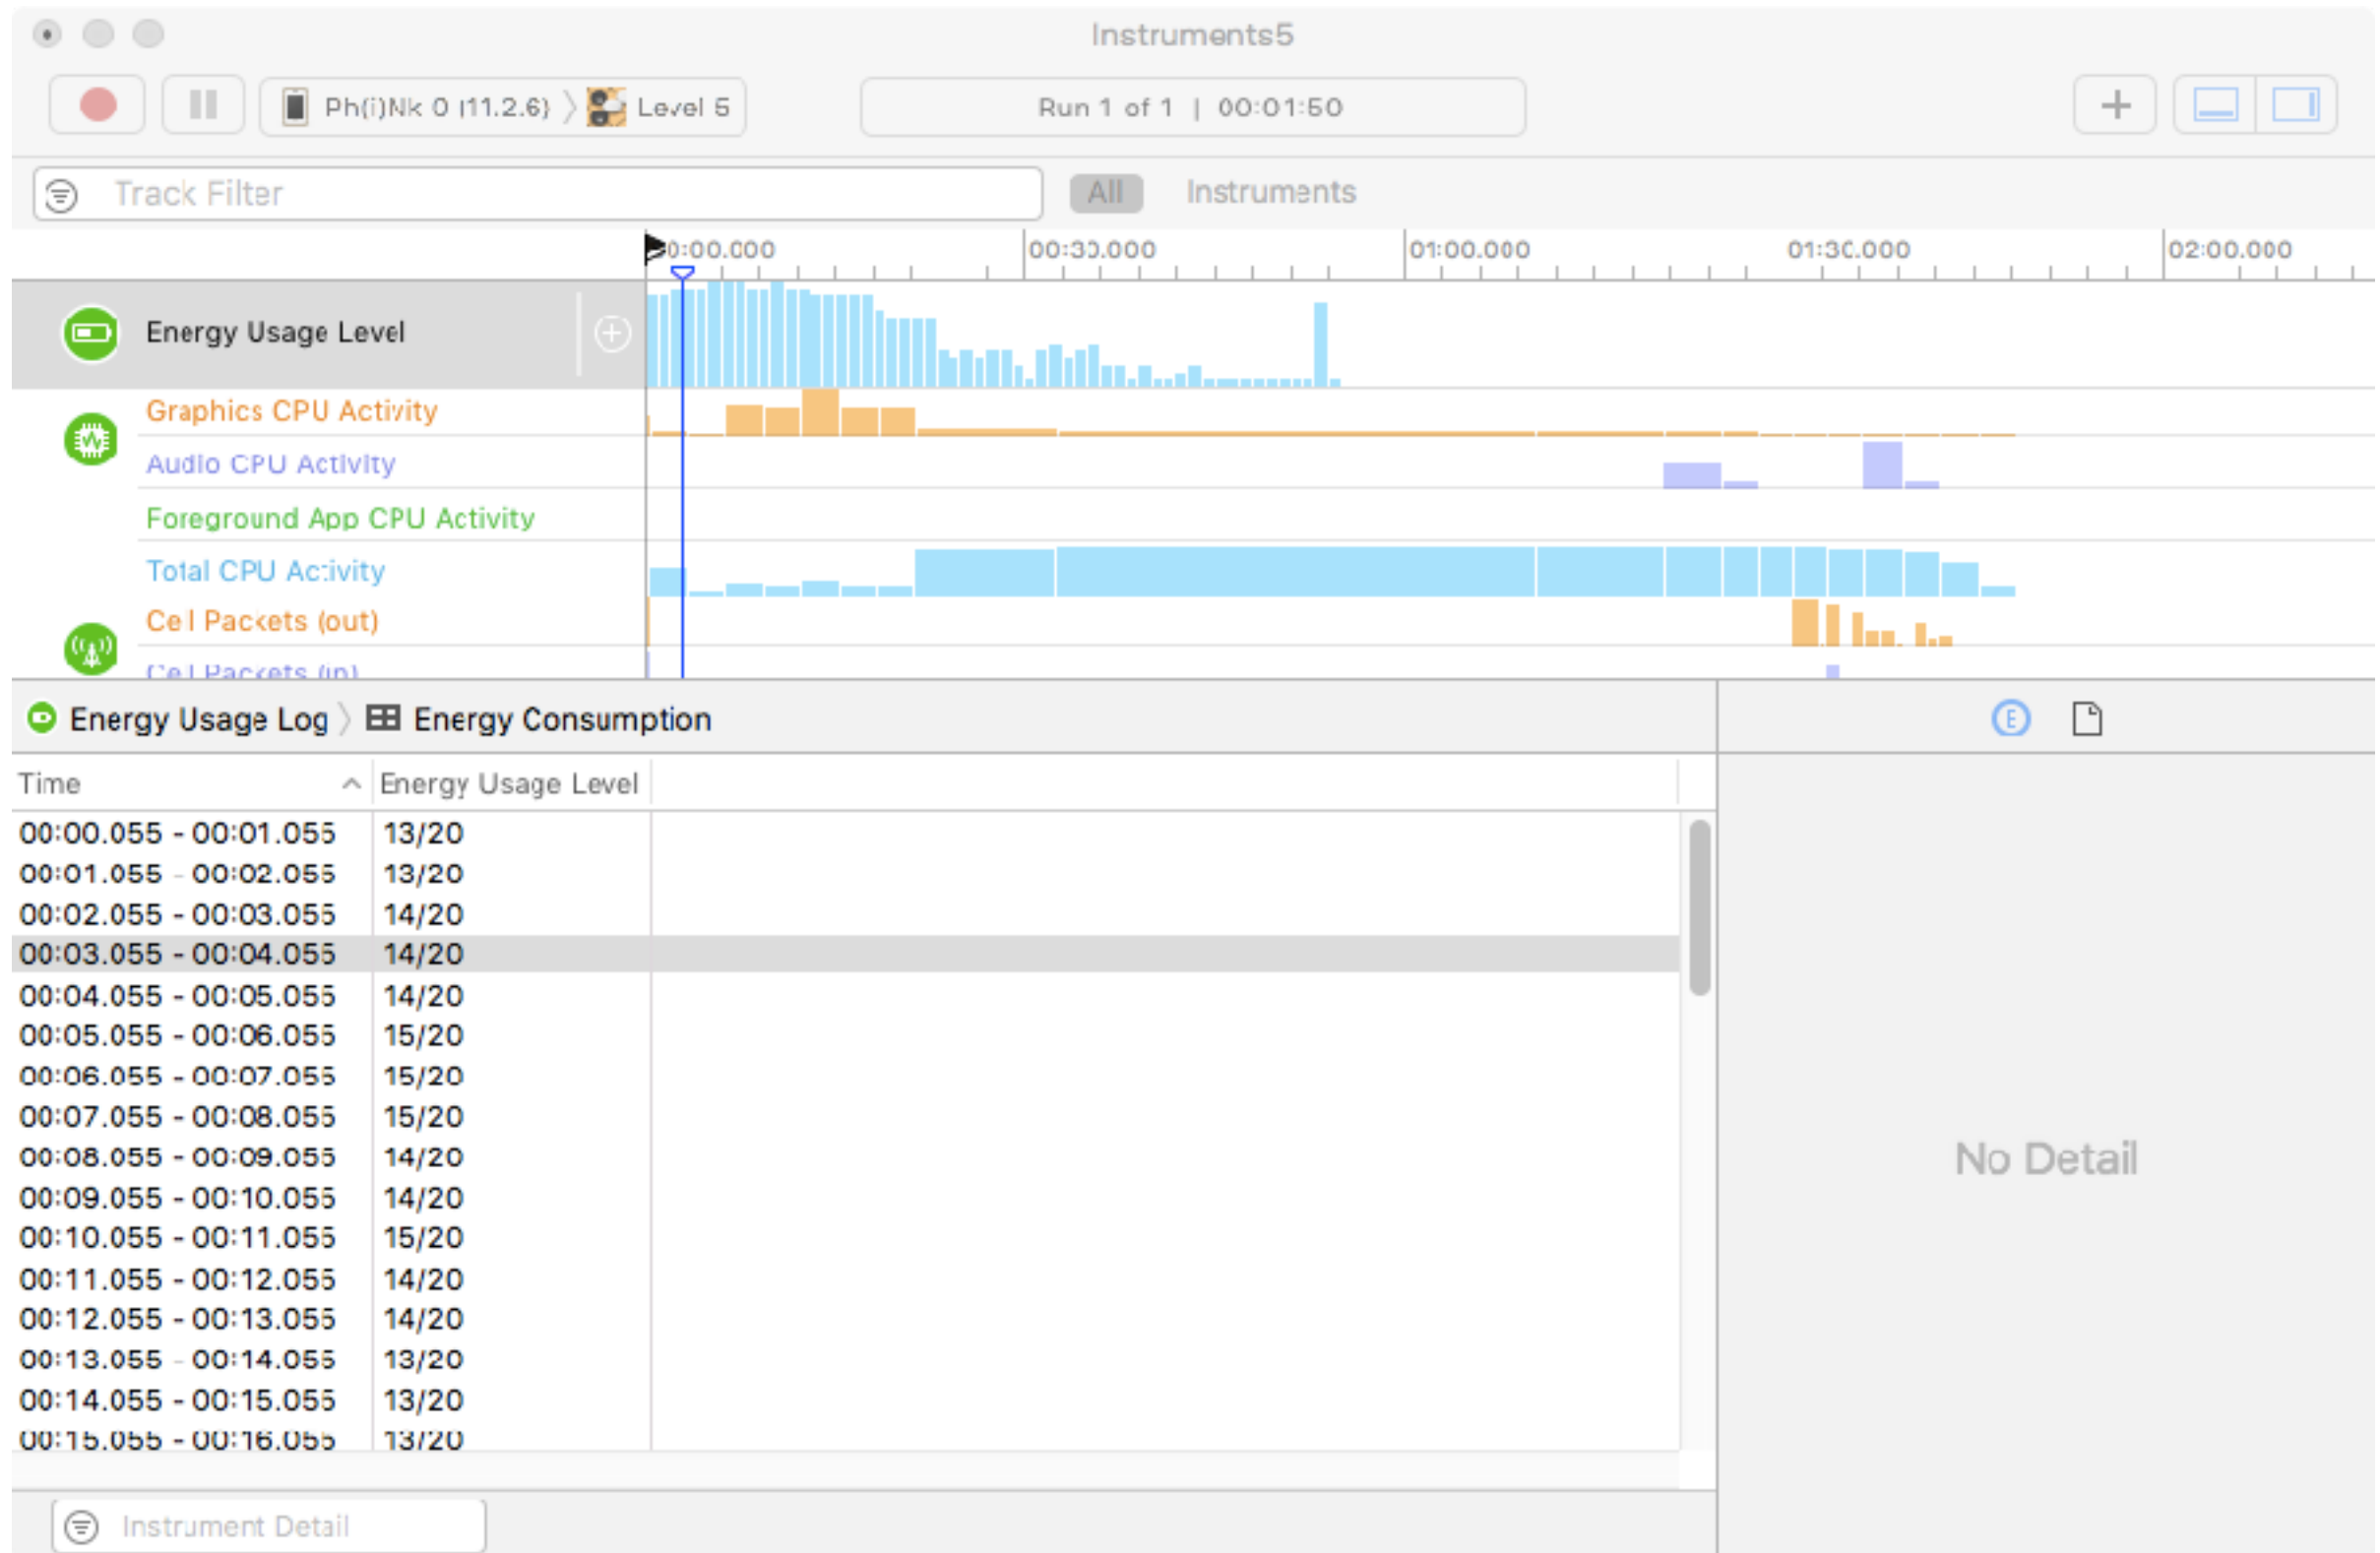


CoreMotion

- Same thing
- Turn sensors off when in background




Energy diagnostics





Measure on the go



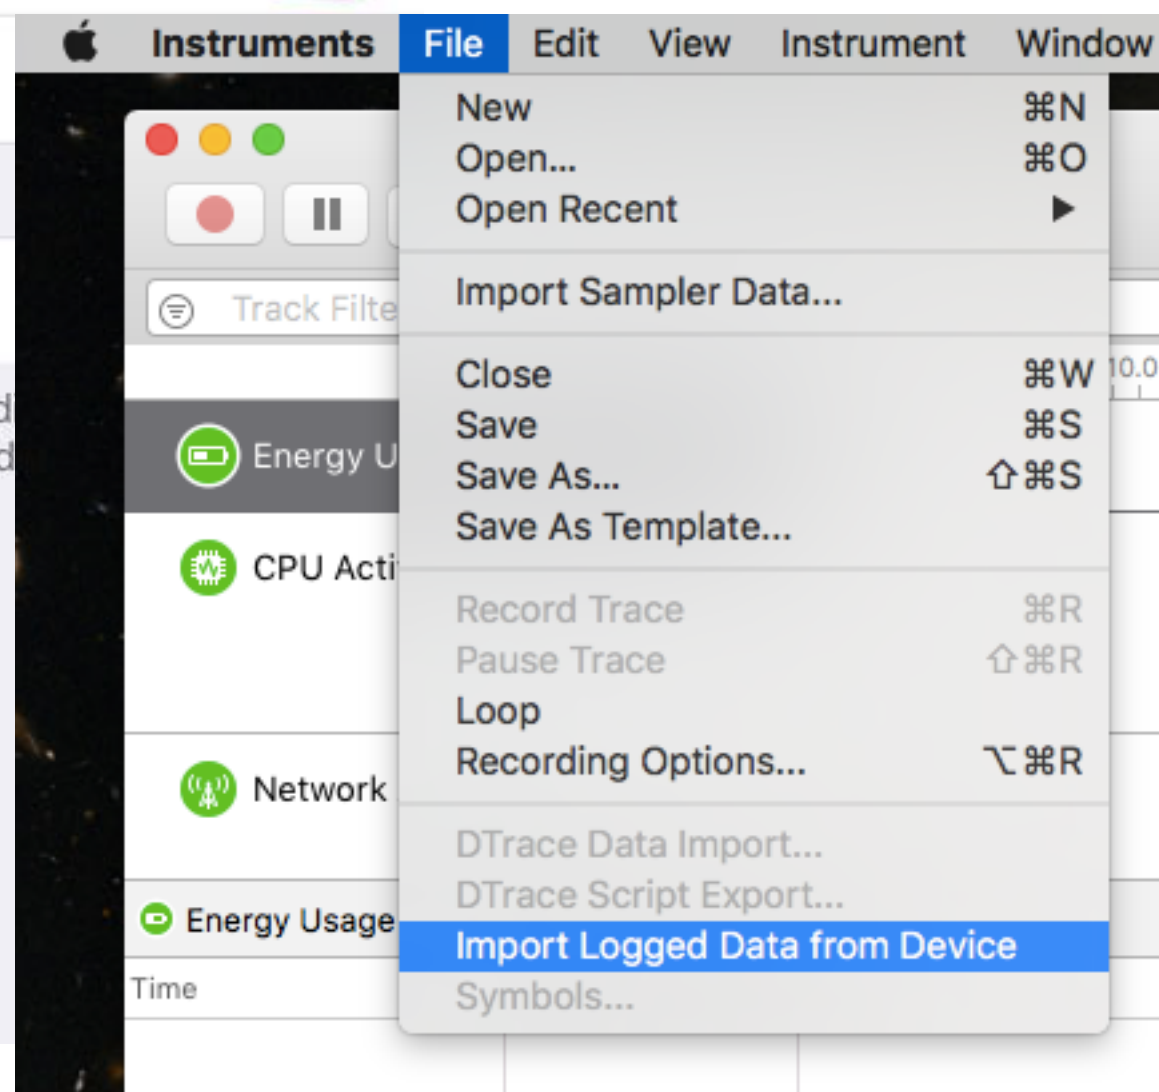
Energy



Networking

Start Recording

Starting a new record
previously recorded d



- Enable device to collect data
- Start/Stop recording
- Import in Instruments

Summary

- Performance optimization is important
- It's also fun
- Remember your data structures, algorithms, compiler classes...
- When? Always!

Resources

Marcel Weiher, *iOS and macOS Performance Tuning: Cocoa, Cocoa Touch, Objective-C, and Swift*. ISBN-13: 978-0321842848

Apple, *Performance Tips*,

<https://developer.apple.com/library/content/documentation/iPhone/Conceptual/iPhoneOSProgrammingGuide/PerformanceTips/PerformanceTips.html>

Apple, *Instruments User Guide*,

https://developer.apple.com/library/content/documentation/DeveloperTools/Conceptual/InstrumentsUserGuide/index.html#//apple_ref/doc/uid/TP40004652

Apple, *Understanding Swift Performance, WWDC 2016*

<https://developer.apple.com/videos/play/wwdc2016/416/>


Writing High-Performance Swift Code,

<https://github.com/apple/swift/blob/master/docs/OptimizationTips.rst>

Performance, optimization and why it matters

Marco Scheurer

Lausanne, Switzerland / Shanghai, China



marco@sente.ch



phink0